

EECS 545: Machine Learning

Lecture 4. Classification

Honglak Lee

1/21/2026





Outline

- Logistic regression (for binary classification)
 - Gradient Descent
 - Newton's method
- Softmax regression (for multi-class classification)
- K-nearest neighbors (KNN)

Supervised learning: classification

Supervised learning

- Goal:
 - Given data X in feature space and labels Y
 - Learn to predict Y from X
- Labels could be discrete or continuous
 - Discrete-valued labels: classification (today's topic)
 - Continuous-valued labels: regression



Classification problem

- The task of classification:
 - Given an input vector \mathbf{x} , assign it to one of K distinct classes C_k where $k = 1, \dots, K$
- Representing the assignment:
 - For $K = 2$:
 - $y = 1$ means that \mathbf{x} is in C_1
 - $y = 0$ means that \mathbf{x} is in C_2 .
 - (Sometimes, $y = -1$ can be used depending on algorithms)
 - For $K > 2$:
 - Use 1-of- K coding
 - e.g., $\mathbf{y} = (0, 1, 0, 0, 0)^T$ means that \mathbf{x} is in C_2 .
 - (This works for $K = 2$ as well)

Classification problem



- Training: train a classifier $h(\mathbf{x})$ from training data
 - Training data $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(N)}, y^{(N)})\}$
- Testing (evaluation):
 - testing data: $h(x_{\text{test}}^{(1)}), h(x_{\text{test}}^{(2)}), \dots, h(x_{\text{test}}^{(N')})$
 - The learning algorithm produces predictions
- 0-1 loss: Classification error = $\frac{1}{N'} \sum_{j=1}^{N'} \mathbb{I} [h(x_{\text{test}}^{(j)}) \neq y_{\text{test}}^{(j)}]$

Logistic regression


Probabilistic discriminative models

- Model decision boundary as a function of input \mathbf{x}
 - Learn $P(C_k | \mathbf{x})$ over data (e.g., maximum likelihood)
 - Directly predict class labels from inputs
- Next class: we will cover probabilistic generative models
 - Learn $P(C_k, \mathbf{x})$ over data (maximum likelihood) and then use Bayes' rule to predict $P(C_k | \mathbf{x})$

Example (1-dim. case)



Example (2-dim. case)



note: we use $\phi(x) = x$
here for illustration.
The decision boundary
can be nonlinear when
 $\phi(x)$ is a nonlinear
function of x .

Logistic regression

- Models the class posterior using a sigmoid applied to a linear function of the feature vector:

$$p(C_1|\phi) = h(\phi) = \sigma(\mathbf{w}^\top \phi(\mathbf{x}))$$

- We can solve the parameter \mathbf{w} by maximizing the likelihood of the training data

Sigmoid and logit functions

- The *logistic sigmoid* function is:


$$\sigma(a) = \frac{1}{1 + \exp(-a)} = \frac{\exp(a)}{1 + \exp(a)}$$

- Its inverse is the *logit* function (aka log odds ratio):

$$a = \ln\left(\frac{\sigma}{1 - \sigma}\right)$$

- Generalizes to *normalized exponential*, or *softmax*

$$p_i = \frac{\exp(q_i)}{\sum_j \exp(q_j)}$$



Class-conditional probability (for a single example)

- Depending on the label y , the conditional probability of y given \mathbf{x} is defined as:

$$P(y = 1|\mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^\top \phi(\mathbf{x}))$$

$$P(y = 0|\mathbf{x}, \mathbf{w}) = 1 - \sigma(\mathbf{w}^\top \phi(\mathbf{x}))$$

- Therefore we can write both cases compactly as:

$$P(y|\mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^\top \phi(\mathbf{x}))^y (1 - \sigma(\mathbf{w}^\top \phi(\mathbf{x})))^{1-y}$$

Likelihood function (of logistic regression)

- The likelihood of Data $\{(\phi(\mathbf{x}^{(n)}), y^{(n)})\}$, where $y^{(n)} \in \{0, 1\}$

$$P(D|\mathbf{w}) = \prod_{i=1}^N P(\mathbf{x}^{(i)}, y^{(i)}|\mathbf{w}) \quad \text{IID (Independent Identical Distribution)}$$

Definition of
conditional
probability

$$\begin{aligned} &= \prod_{i=1}^N P(y^{(i)}|\mathbf{x}^{(i)}, \mathbf{w}) \underbrace{P(\mathbf{x}^{(i)}|\mathbf{w})}_{=P(\mathbf{x}^{(i)})} \\ &\qquad\qquad\qquad \boxed{\text{P(x) does not depend on w}} \end{aligned}$$

$$\propto \prod_{i=1}^N P(y^{(i)}|\mathbf{x}^{(i)}, \mathbf{w}) \longrightarrow P(\mathbf{y}|\mathbf{X}, \mathbf{w})$$

Compact notation: Technically speaking, this is
(conditional) likelihood of y given X

Logistic regression

- For a data set $\{(\phi(\mathbf{x}^{(n)}), y^{(n)})\}$, where $y^{(n)} \in \{0, 1\}$ the likelihood function is

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N (h^{(n)})^{y^{(n)}} (1 - h^{(n)})^{1-y^{(n)}}$$

where

$$h^{(n)} = p(C_1 | \phi(\mathbf{x}^{(n)})) = \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}))$$

- Define a loss function $E(\mathbf{w}) = -\log p(\mathbf{y}|\mathbf{X}, \mathbf{w})$
 - Minimizing $E(\mathbf{w})$ maximizes likelihood

note: $h(\mathbf{x})$ is the hypothesis function,
 $\sigma(\mathbf{x})$ is the specific hypothesis for logistic regression

Derivation

- $\log P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N y^{(n)} \log h^{(n)} + (1 - y^{(n)}) \log(1 - h^{(n)})$
- Gradient (matrix calculus)

$$\begin{aligned}\nabla_{\mathbf{w}} \log P(\mathbf{y}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}, \mathbf{w}) \\ = \sum_{n=1}^N \nabla_{\mathbf{w}} \left(y^{(n)} \log h(\mathbf{x}^{(n)}, \mathbf{w}) + (1 - y^{(n)}) \log(1 - h(\mathbf{x}^{(n)}, \mathbf{w})) \right)\end{aligned}$$

Derivation

- $\log P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N y^{(n)} \log h^{(n)} + (1 - y^{(n)}) \log(1 - h^{(n)})$
- Gradient (matrix calculus) $h(\mathbf{x}^{(n)}, \mathbf{w}) \triangleq \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \triangleq \sigma^{(n)}$

$$\begin{aligned}\nabla_{\mathbf{w}} \log P(\mathbf{y}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}, \mathbf{w}) \\ = \sum_{n=1}^N \nabla_{\mathbf{w}} \left(y^{(n)} \log h(\mathbf{x}^{(n)}, \mathbf{w}) + (1 - y^{(n)}) \log(1 - h(\mathbf{x}^{(n)}, \mathbf{w})) \right)\end{aligned}$$


Derivation

- $\log P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N y^{(n)} \log h^{(n)} + (1 - y^{(n)}) \log(1 - h^{(n)})$
- Gradient (matrix calculus) $h(\mathbf{x}^{(n)}, \mathbf{w}) \triangleq \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \triangleq \sigma^{(n)}$

$$\nabla_{\mathbf{w}} \log P(\mathbf{y}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}, \mathbf{w})$$

$$= \sum_{n=1}^N \nabla_{\mathbf{w}} \left(y^{(n)} \log h(\mathbf{x}^{(n)}, \mathbf{w}) + (1 - y^{(n)}) \log(1 - h(\mathbf{x}^{(n)}, \mathbf{w})) \right)$$

$$= \sum_{n=1}^N \left(y^{(n)} \frac{\sigma^{(n)}(1 - \sigma^{(n)})}{\sigma^{(n)}} - (1 - y^{(n)}) \frac{\sigma^{(n)}(1 - \sigma^{(n)})}{1 - \sigma^{(n)}} \right) \nabla_{\mathbf{w}} (\mathbf{w}^\top \phi(\mathbf{x}^{(n)}))$$

$$\boxed{\frac{\partial}{\partial s} \sigma(s) = \frac{\partial}{\partial s} \left(\frac{1}{1 + \exp(-s)} \right) = \sigma(s)(1 - \sigma(s))}$$



Derivation

- $\log P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N y^{(n)} \log h^{(n)} + (1 - y^{(n)}) \log(1 - h^{(n)})$
- **Gradient (matrix calculus)** $h(\mathbf{x}^{(n)}, \mathbf{w}) \triangleq \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \triangleq \sigma^{(n)}$
$$\begin{aligned} & \nabla_{\mathbf{w}} \log P(\mathbf{y}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}, \mathbf{w}) \\ &= \sum_{n=1}^N \nabla_{\mathbf{w}} \left(y^{(n)} \log h(\mathbf{x}^{(n)}, \mathbf{w}) + (1 - y^{(n)}) \log(1 - h(\mathbf{x}^{(n)}, \mathbf{w})) \right) \\ &= \sum_{n=1}^N \left(y^{(n)} \frac{\sigma^{(n)}(1 - \sigma^{(n)})}{\sigma^{(n)}} - (1 - y^{(n)}) \frac{\sigma^{(n)}(1 - \sigma^{(n)})}{1 - \sigma^{(n)}} \right) \nabla_{\mathbf{w}} (\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \\ &= \sum_{n=1}^N \left(y^{(n)}(1 - \sigma^{(n)}) - (1 - y^{(n)})\sigma^{(n)} \right) \nabla_{\mathbf{w}} (\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \end{aligned}$$

Derivation

- $\log P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N y^{(n)} \log h^{(n)} + (1 - y^{(n)}) \log(1 - h^{(n)})$
- **Gradient (matrix calculus)** $h(\mathbf{x}^{(n)}, \mathbf{w}) \triangleq \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \triangleq \sigma^{(n)}$

$$\begin{aligned}\nabla_{\mathbf{w}} \log P(\mathbf{y}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}, \mathbf{w}) \\ &= \sum_{n=1}^N \nabla_{\mathbf{w}} \left(y^{(n)} \log h(\mathbf{x}^{(n)}, \mathbf{w}) + (1 - y^{(n)}) \log(1 - h(\mathbf{x}^{(n)}, \mathbf{w})) \right) \\ &= \sum_{n=1}^N \left(y^{(n)} \frac{\sigma^{(n)}(1 - \sigma^{(n)})}{\sigma^{(n)}} - (1 - y^{(n)}) \frac{\sigma^{(n)}(1 - \sigma^{(n)})}{1 - \sigma^{(n)}} \right) \nabla_{\mathbf{w}} (\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \\ &= \sum_{n=1}^N \left(y^{(n)}(1 - \sigma^{(n)}) - (1 - y^{(n)})\sigma^{(n)} \right) \nabla_{\mathbf{w}} (\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) \\ &= \sum_{n=1}^N \left(y^{(n)} - \sigma^{(n)} \right) \phi(\mathbf{x}^{(n)})\end{aligned}$$

Logistic regression: gradient descent

Note: Vectorized form

- Taking the gradient of $E(\mathbf{w})$ gives us

$$\nabla E(\mathbf{w}) = \Phi^T (\mathbf{h} - \mathbf{y})$$

$$\nabla_{\mathbf{w}} E(\mathbf{w}) = -\nabla_{\mathbf{w}} \log P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N (h^{(n)} - y^{(n)}) \phi(\mathbf{x}^{(n)})$$

- Recall $h^{(n)} = p(C_1 | \phi(\mathbf{x}^{(n)})) = \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)}))$
- This is essentially the same gradient expression that appeared in linear regression with least-squares.
- Note the error term between model prediction and target value:
 - Logistic regression: $h^{(n)} - y^{(n)} = \sigma(\mathbf{w}^\top \phi(\mathbf{x}^{(n)})) - y^{(n)}$
 - Cf. Linear regression: $h^{(n)} - y^{(n)} = \mathbf{w}^\top \phi(\mathbf{x}^{(n)}) - y^{(n)}$

Note: Gradient Descent vs Ascent

The following optimization problems are equivalent:

Minimizing the negative log-likelihood

→ gradient **descent** with $E(\mathbf{w}) = -\log P(\mathbf{y}|\mathbf{X}, \mathbf{w})$ as an objective

$$w := w - \eta \nabla_{\mathbf{w}} E(\mathbf{w}) = w - \eta \sum_{n=1}^N (h^{(n)} - y^{(n)}) \phi(\mathbf{x}^{(n)})$$

Maximizing the log-likelihood

→ gradient **ascent** with $\log P(\mathbf{y}|\mathbf{X}, \mathbf{w})$ as an objective

$$w := w + \eta \nabla_{\mathbf{w}} \log P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = w + \eta \sum_{n=1}^N (y^{(n)} - h^{(n)}) \phi(\mathbf{x}^{(n)})$$

Newton's method

- Goal: Minimizing a general function $E(\mathbf{w})$ (one-dimensional case)

- Approach: solve for

$$f(\mathbf{w}) = \frac{\partial E(\mathbf{w})}{\partial \mathbf{w}} = 0$$




- So, how to solve this problem?

- Newton's method (aka Newton-Raphson method)
 - Repeat until convergence:

$$\mathbf{w} := \mathbf{w} - \frac{f(\mathbf{w})}{f'(\mathbf{w})}$$

Newton's method

- Interactively solve until we get $f(w) = 0$.



- Geometric intuition:

$$w := w - \frac{f(w)}{f'(w)}$$

Current value
“Slope”

Newton's method

- Now we want to minimize $E(\mathbf{w})$
 - Convert $E'(\mathbf{w}) = f(\mathbf{w})$
 - Repeat until convergence

$$\mathbf{w} := \mathbf{w} - \frac{E'(\mathbf{w})}{E''(\mathbf{w})}$$

Newton update
when \mathbf{w} is a scalar

Newton's method

- Now we want to minimize $E(\mathbf{w})$
 - Convert $E'(\mathbf{w}) = f(\mathbf{w})$
 - Repeat until convergence

$$\mathbf{w} := \mathbf{w} - \frac{E'(\mathbf{w})}{E''(\mathbf{w})}$$

Newton update
when \mathbf{w} is a scalar

- This method can be extended to the multivariate case:

$$\mathbf{w} := \mathbf{w} - H^{-1} \nabla_{\mathbf{w}} E$$

Newton update
when \mathbf{w} is a vector

where H is a Hessian matrix evaluated at \mathbf{w}

$$H_{ij}(\mathbf{w}) = \frac{\partial^2 E(\mathbf{w})}{\partial \mathbf{w}_i \partial \mathbf{w}_j}$$

- Note: for linear regression, the Hessian is $\Phi^\top \Phi$

Derivation of Newton's method

Taylor expansion of $E(\mathbf{w})$ at \mathbf{w}_0 up to 2nd order:

$$\begin{aligned} E(\mathbf{w}) \approx & E(\mathbf{w}_0) + \nabla E(\mathbf{w}_0)^T (\mathbf{w} - \mathbf{w}_0) \\ & + \frac{1}{2} (\mathbf{w} - \mathbf{w}_0)^T H(\mathbf{w}_0) (\mathbf{w} - \mathbf{w}_0) \end{aligned}$$

Find a closed-form solution that optimizes the quadratic approximation above

$$\nabla [E(\mathbf{w}_0) + \nabla E(\mathbf{w}_0)^T (\mathbf{w} - \mathbf{w}_0) + \frac{1}{2} (\mathbf{w} - \mathbf{w}_0)^T H(\mathbf{w}_0) (\mathbf{w} - \mathbf{w}_0)] = \mathbf{0}$$

Using matrix calculus trick similar to Linear regression, we get

$$\mathbf{w} = \mathbf{w}_0 - H(\mathbf{w}_0)^{-1} \nabla E(\mathbf{w}_0)$$

Iteratively Reweighted Least Squares (IRLS)

- Recall: for linear regression, least-squares has a closed-form solution:

$$\mathbf{w}_{\text{ML}} = (\Phi^\top \Phi)^{-1} \Phi^\top \mathbf{y}$$

- This generalizes to weighted-least-squares with an NxN diagonal weight matrix \mathbf{R} .

$$\mathbf{w}_{\text{WLS}} = (\Phi^\top \mathbf{R} \Phi)^{-1} \Phi^\top \mathbf{R} \mathbf{y}$$

- For logistic regression, however, $h(\mathbf{x}, \mathbf{w})$ is non-linear, and there is no closed-form solution. So **we need to iterate (i.e. repeatedly apply Newton steps) to get the optimal solution, which is called IRLS.**

Iterative solution

- Apply Newton-Raphson method to iterate to a solution \mathbf{w} for $\nabla E(\mathbf{w}) = 0$
- This involves least-squares with weights \mathbf{R} :

$$R_{nn} = h^{(n)}(1 - h^{(n)})$$

- Since \mathbf{R} depends on \mathbf{w} (and vice versa), we get *iterative reweighted least squares* (IRLS)

$$\text{where } \mathbf{w}^{(\text{new})} = (\Phi^\top \mathbf{R} \Phi)^{-1} \Phi^\top \mathbf{R} \mathbf{z}$$

$$\mathbf{z} = \Phi \mathbf{w}^{(\text{old})} - \mathbf{R}^{-1}(\mathbf{h} - \mathbf{y})$$

Iterative solution: Derivation

Applying Newton's method:

$$\mathbf{w}^{(\text{new})} = \mathbf{w}^{(\text{old})} - H(\mathbf{w}^{(\text{old})})^{-1} \nabla E(\mathbf{w}^{(\text{old})})$$

Gradient and Hessians from Logistic Regression loss function:

$$\nabla E(\mathbf{w}) = \Phi^T (\mathbf{h} - \mathbf{y}) \quad \text{and} \quad H(\mathbf{w}) = \nabla^2 E(\mathbf{w}) = \Phi^T R \Phi$$

where $R_{nn} = h^{(n)} (1 - h^{(n)})$

Putting together: $\mathbf{w}^{(\text{new})} = \mathbf{w}^{(\text{old})} - [\Phi^T R \Phi]^{-1} \Phi^T (\mathbf{h} - \mathbf{y})$

If we define \mathbf{z} as $\mathbf{z} = \Phi \mathbf{w}^{(\text{old})} - R^{-1} (\mathbf{h} - \mathbf{y})$

We can also write the solution as: $\mathbf{w}^{(\text{new})} = (\Phi^T R \Phi)^{-1} \Phi^T R \mathbf{z}$

Multi-class Classification

Softmax regression for multiclass classification

- For multiclass case, we can use softmax regression.
 - Softmax regression can be viewed as a generalization of logistic regression
- Recall that, logistic regression (binary classification) models class conditional probability as:

$$p(y = 1 \mid \mathbf{x}; \mathbf{w}) = \frac{\exp(\mathbf{w}^\top \phi(\mathbf{x}))}{1 + \exp(\mathbf{w}^\top \phi(\mathbf{x}))}$$

$$p(y = 0 \mid \mathbf{x}; \mathbf{w}) = \frac{1}{1 + \exp(\mathbf{w}^\top \phi(\mathbf{x}))}$$

- Note that these probabilities sum to 1.



note: we use $\phi(\mathbf{x}) = \mathbf{x}$ here for illustration.
The decision boundary can be nonlinear when $\phi(\mathbf{x})$ is a nonlinear function of \mathbf{x} .

- For multiclass classification (with K classes), we use the following model

$$p(y = k \mid \mathbf{x}; \mathbf{w}) = \frac{\exp(\mathbf{w}_k^\top \phi(\mathbf{x}))}{1 + \sum_{j=1}^{K-1} \exp(\mathbf{w}_j^\top \phi(\mathbf{x}))} \quad \text{for } k = \{1, \dots, K-1\}$$

$$p(y = K \mid \mathbf{x}; \mathbf{w}) = \frac{1}{1 + \sum_{j=1}^{K-1} \exp(\mathbf{w}_j^\top \phi(\mathbf{x}))} \quad \text{equivalent to setting } \mathbf{w}_K = 0$$

- Note that these probabilities sum to 1.



Softmax regression: Log-likelihood (objective function) and learning

- Defining $w_K = 0$, we can write as:

$$p(y = k \mid \mathbf{x}; \mathbf{w}) = \frac{\exp(\mathbf{w}_k^\top \phi(\mathbf{x}))}{\sum_{j=1}^K \exp(\mathbf{w}_j^\top \phi(\mathbf{x}))}$$

$$p(y \mid \mathbf{x}; \mathbf{w}) = \prod_{k=1}^K \left[\frac{\exp(\mathbf{w}_k^\top \phi(\mathbf{x}))}{\sum_{j=1}^K \exp(\mathbf{w}_j^\top \phi(\mathbf{x}))} \right]^{\mathbb{I}(y=k)}$$



- Log-Likelihood $\log p(D|\mathbf{w}) = \sum_i \log p(y^{(i)} | \mathbf{x}^{(i)}, \mathbf{w})$
 $= \sum_i \log \prod_{k=1}^K \left[\frac{\exp(\mathbf{w}_k^\top \phi(\mathbf{x}^{(i)}))}{\sum_{j=1}^K \exp(\mathbf{w}_j^\top \phi(\mathbf{x}^{(i)}))} \right]^{\mathbb{I}(y^{(i)}=k)}$
- We can learn \mathbf{w} by gradient ascent for maximizing the log-likelihood or iterative Newton's method (IRLS).

K-nearest neighbor classification

K-nearest neighbors


- Training method:
 - Save the training examples (no sophisticated learning)
- At prediction (testing) time:
 - Given a test (query) example \mathbf{x} , find the K training examples that are *closest* to \mathbf{x} .
$$\text{KNN}(\mathbf{x}) = \{ (\mathbf{x}^{(1)\prime}, y^{(1)\prime}), (\mathbf{x}^{(2)\prime}, y^{(2)\prime}), \dots, (\mathbf{x}^{(K)\prime}, y^{(K)\prime}) \}$$
- Predict the most frequent class among all y 's from $\text{KNN}(\mathbf{x})$.
$$h(\mathbf{x}) = \arg \max_y \sum_{(\mathbf{x}', y') \in \text{KNN}(\mathbf{x})} \mathbb{I}[y' = y] \quad \text{"majority vote"}$$
- Note: this function can be applied to regression!

K-nearest neighbors for classification



Slide credit: Ben Kuipers

K-nearest neighbors for classification




- Larger K leads to a smoother decision boundary (bias-variance trade-off)
- Classification performance generally improves as N (training set size) increases
- For $N \Rightarrow \infty$, the error rate of the 1-nearest-neighbor classifier is never more than twice the optimal error (obtained from the true conditional class distributions). See ESL CH 13.3.

Factors (hyperparameters) affecting KNN

- Distance metric $D(\mathbf{x}, \mathbf{x}')$
 - How to define distance between two examples \mathbf{x} and \mathbf{x}' ?
- The value of K
 - K determines how much we “smooth out” the prediction

What is the decision boundary?

Voronoi diagram: Euclidean (L_2) distance









Note: Each region corresponds kNN's prediction when K=1

i.e. prediction is the same as the corresponding training sample's label in each region (training sample is visualized as dot).

Dependence on distance metric (L^q norm)

Distance between i-th and j-th example: $\sqrt[q]{\sum_l (x_l^{(i)} - x_l^{(j)})^q}$



KNN: classification vs regression

- We can formulate KNN into regression/classification
- For classification, where the label y is categorical, we take the “majority vote” over target labels.

$$h(\mathbf{x}) = \arg \max_y \sum_{(\mathbf{x}', y') \in \text{KNN}(\mathbf{x})} \mathbb{I}[y' = y]$$

- For regression, where the label y is real-valued numbers, we take “average” over target labels.

$$h(\mathbf{x}) = \frac{1}{k} \sum_{(\mathbf{x}', y') \in \text{KNN}(\mathbf{x})} y'$$

Advantage/disadvantages of KNN methods

- Advantage:
 - Very simple and flexible (no assumption on distribution)
 - Effective (e.g. for low dimensional inputs)
- Disadvantages:
 - Expensive: need to remember (store) and search through all the training data for every prediction
 - Curse of dimensionality: in high dimensions, all points are far
 - Not robust to irrelevant features: if \mathbf{x} has irrelevant/noisy features, then distance function does not reflect similarity between examples

Concept check

- How are labels represented in multiclass classification problems?
- What is the motivation for using Newton's method for optimization in logistic regression?
- What does increasing K do for the results from KNN?

Any feedback (about lecture, slide, homework, project, etc.)?

(via anonymous google form: <https://forms.gle/fpYmiBtG9Me5qbP37>)



Change Log of lecture slides:

<https://docs.google.com/document/d/e/2PACX-1vSSIHjklypK7rKFSR1-5GYXyBCEW8UPtpSfCR9AR6M1l7K9ZQEmxfFwaWaW7kLDxusthsF8WICyZJ-/pub>